Mastering RAG

A Perfect Guide to
Retrieval Augmented Generation

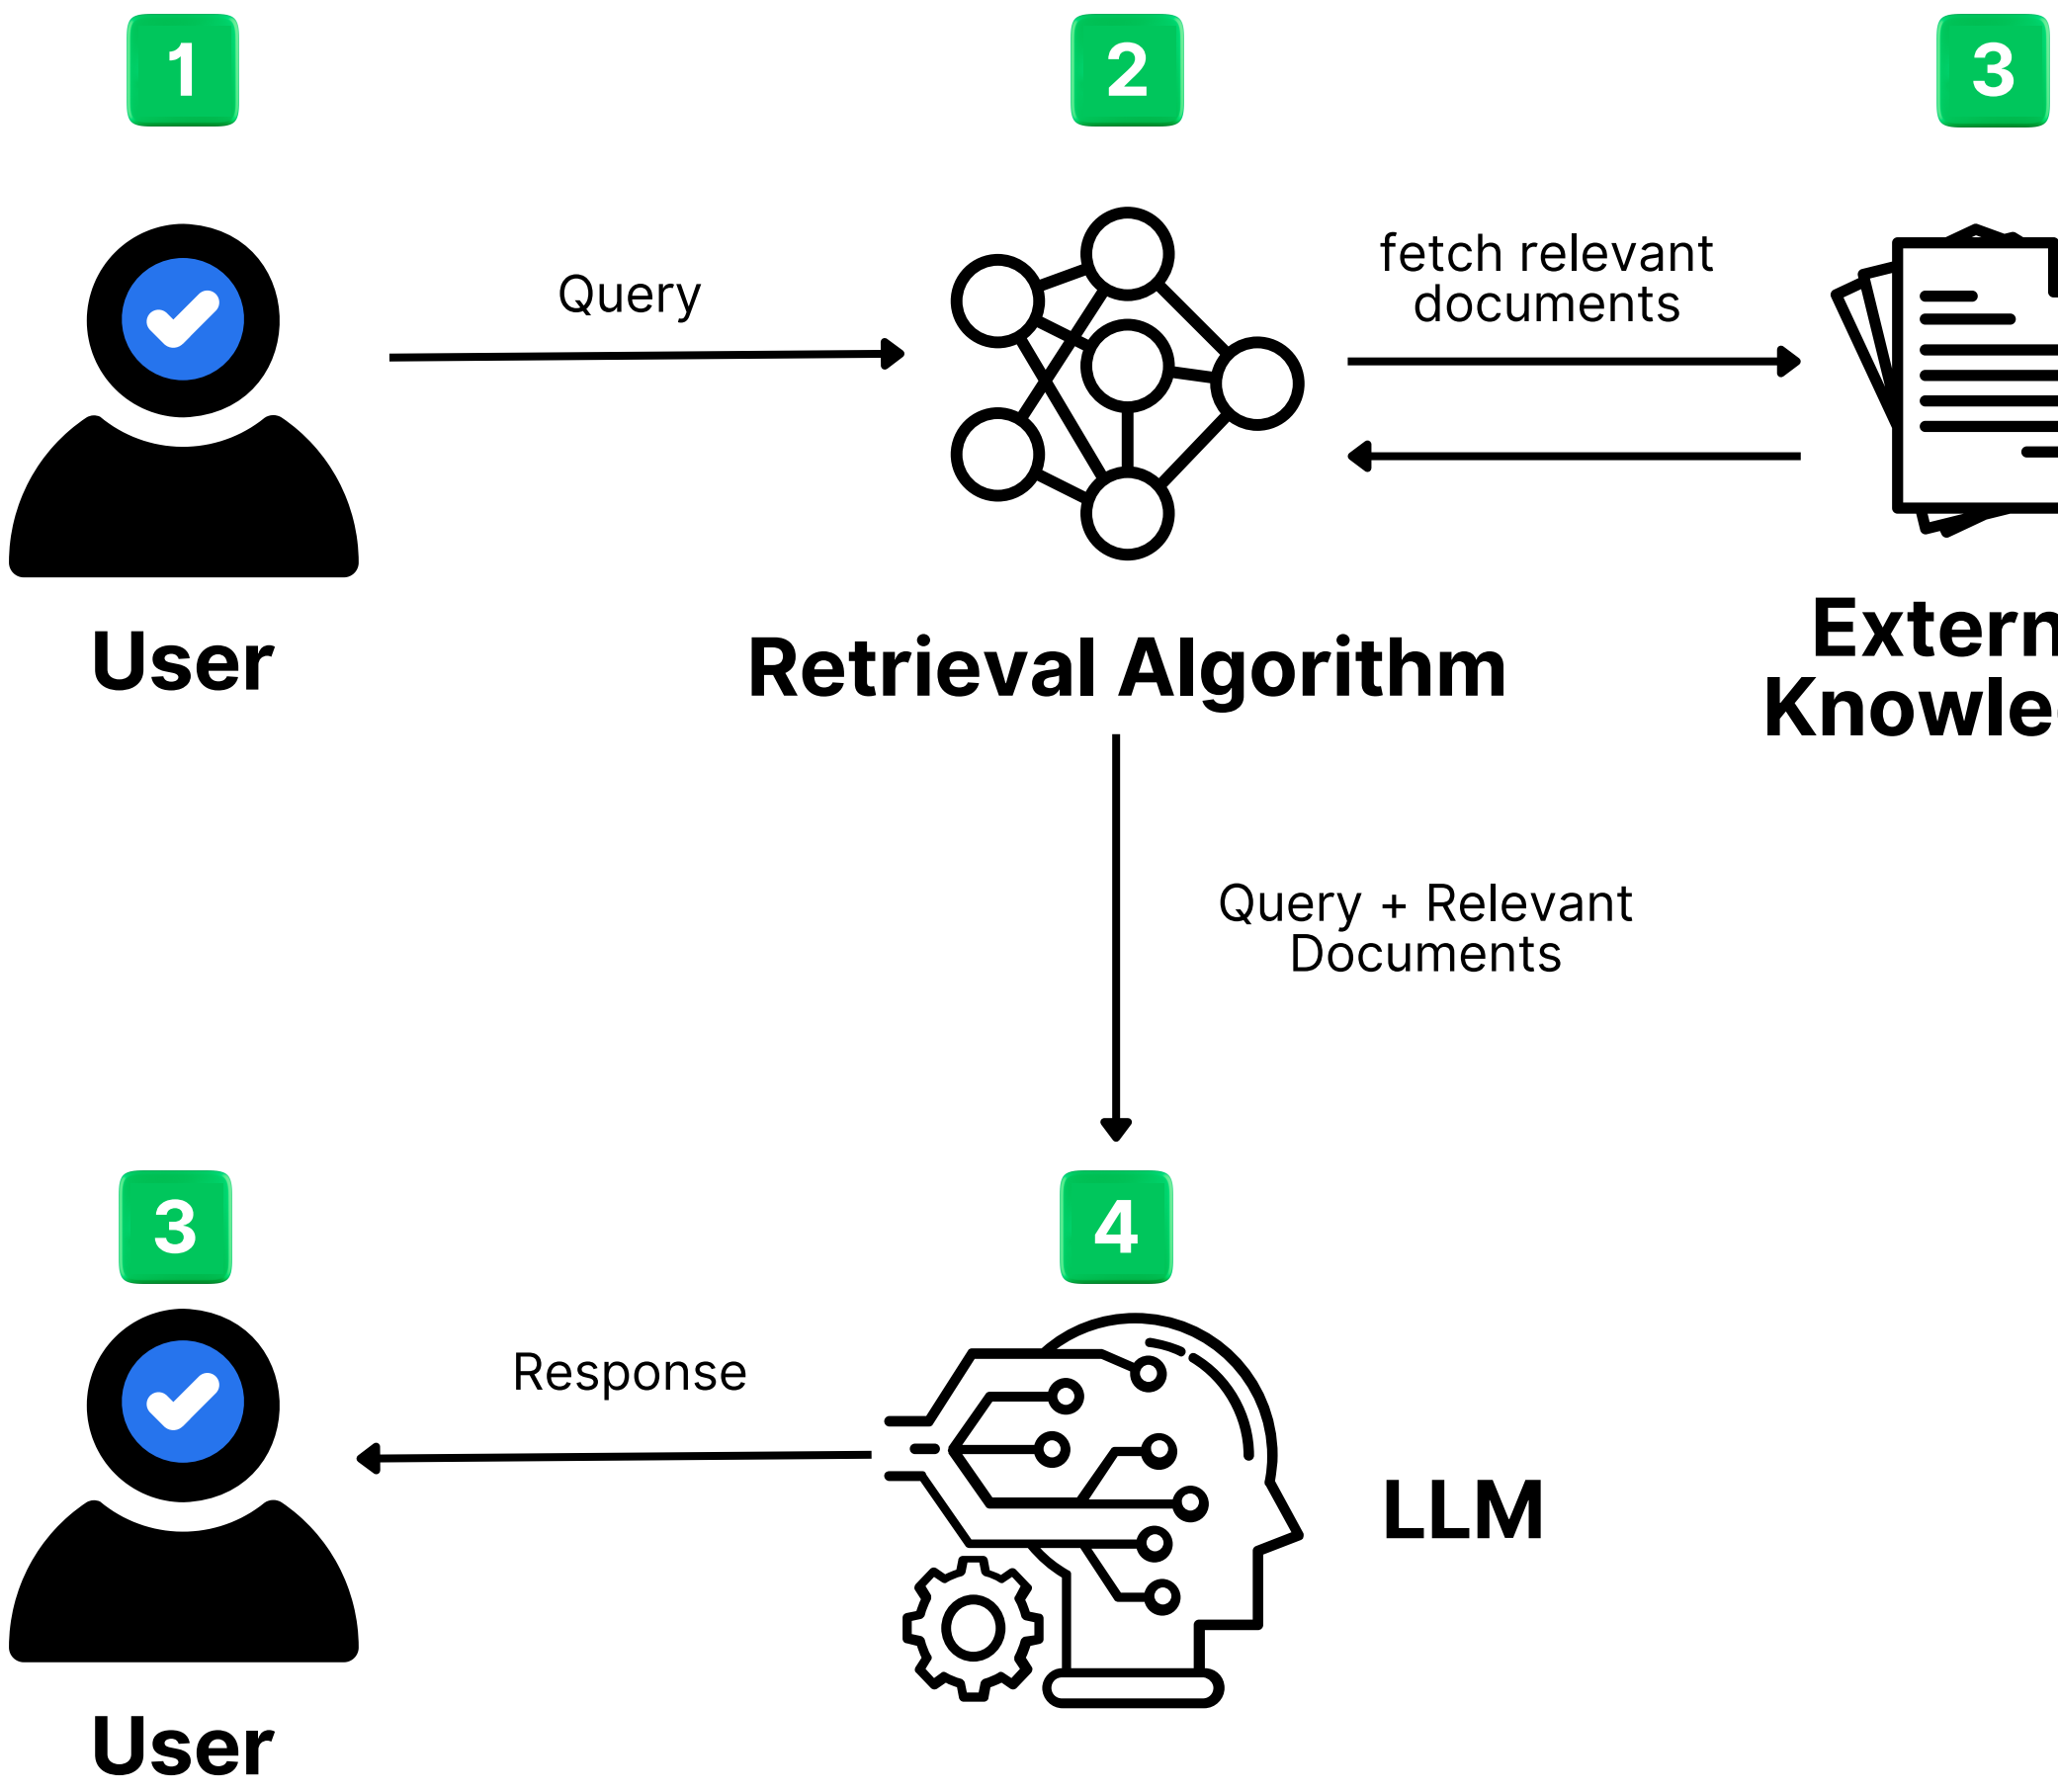

Analytics Vidhya

While building LLMs using Prompt Engineering, there are the following problems that occurs:

- **The model can be inconsistent.**
- **The model can hallucinate.**
- **The model is not up-to date.**

To overcome all these problems, we use one of popular method to build applications powered by LLMs which is

RAG

Ingest custom knowledge to LLMs

What is RAG?

- Retrieval-Augmented Generation (RAG) is an innovative approach in natural language processing that combines two primary components: a retrieval mechanism and a generative model.

- The retrieval component searches a large database of documents to find relevant information, which the generative model then uses to produce a coherent and contextually appropriate response.

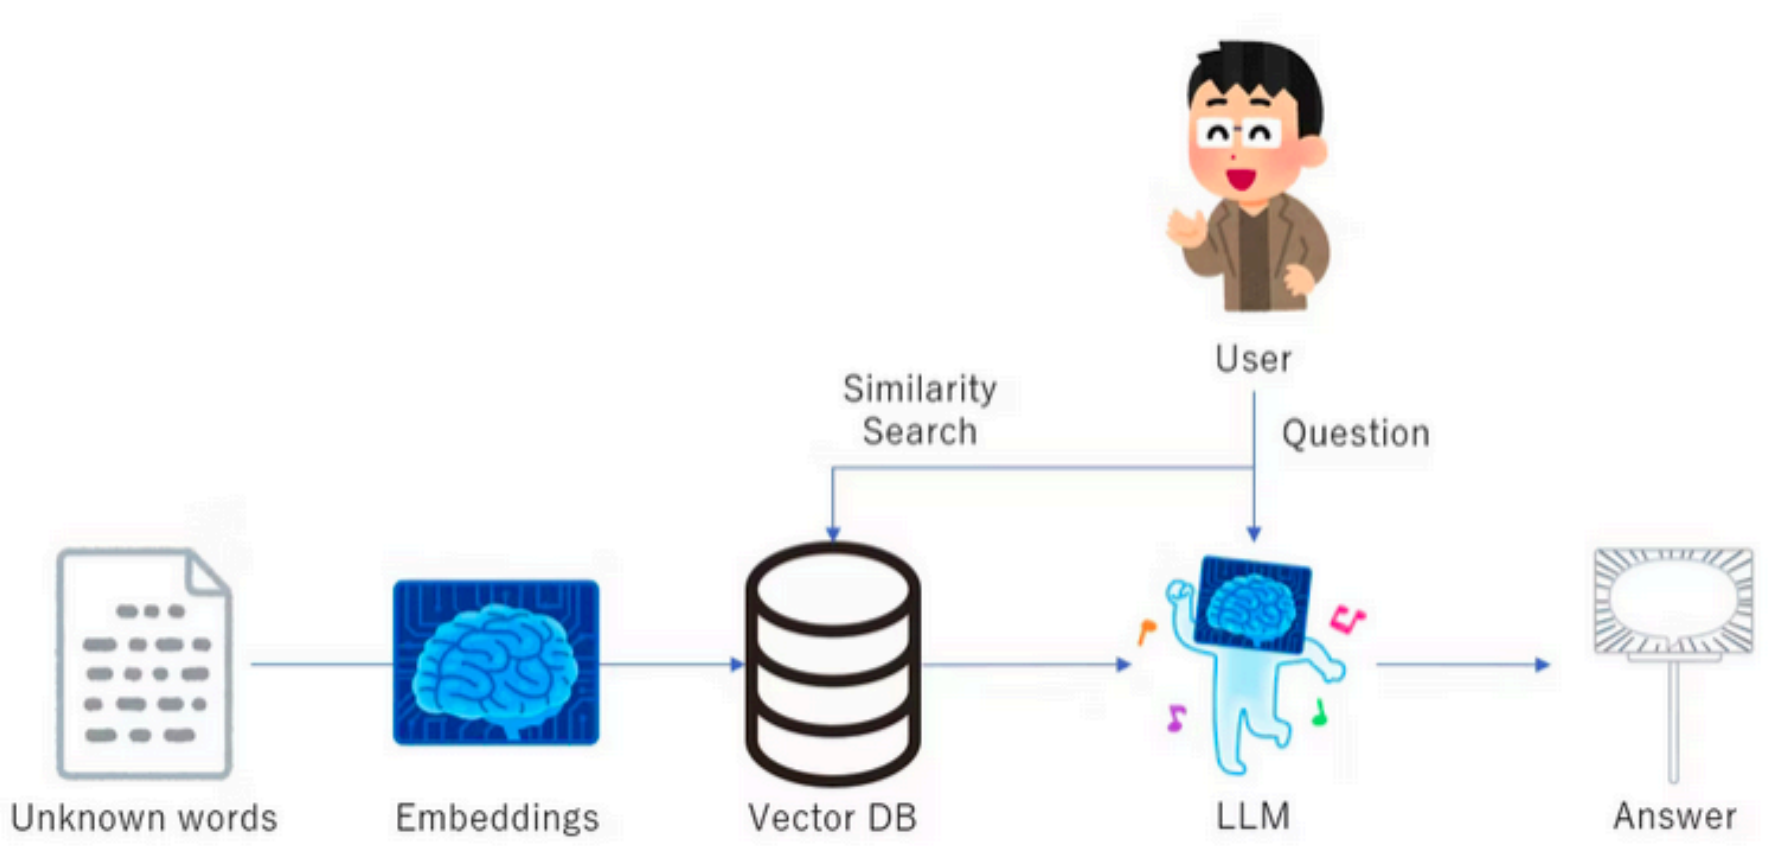

How RAG Differs from LLMs?

- Traditional language models (LLMs) generate responses based solely on their training data and the input query.
- While they can be remarkably effective, they often struggle with providing up-to-date or specific information not present in their training data.
- On the other hand, RAG systems augment their generative capabilities with real-time retrieval of information, ensuring responses are fluent, factually grounded, and relevant.

The Importance of RAG

RAG systems are particularly useful in scenarios where up-to-date and specific information is crucial. Some notable applications include:

- **Customer Support**: Providing accurate and timely responses to customer queries by retrieving relevant information from a knowledge base.

- **Healthcare**: Assisting medical professionals with quick access to the latest research and clinical guidelines.

- **Education**: Offering detailed explanations and additional resources to students based on their queries.

Advantages

- RAG systems can generate more accurate and contextually appropriate responses by retrieving relevant documents.
- These systems can access the latest information, making them ideal for dynamic fields where knowledge constantly evolves.
- The retrieval mechanism ensures the generated responses are closely aligned with the user’s query, enhancing the overall user experience.

Component of RAGs

There are basically two component in RAG:

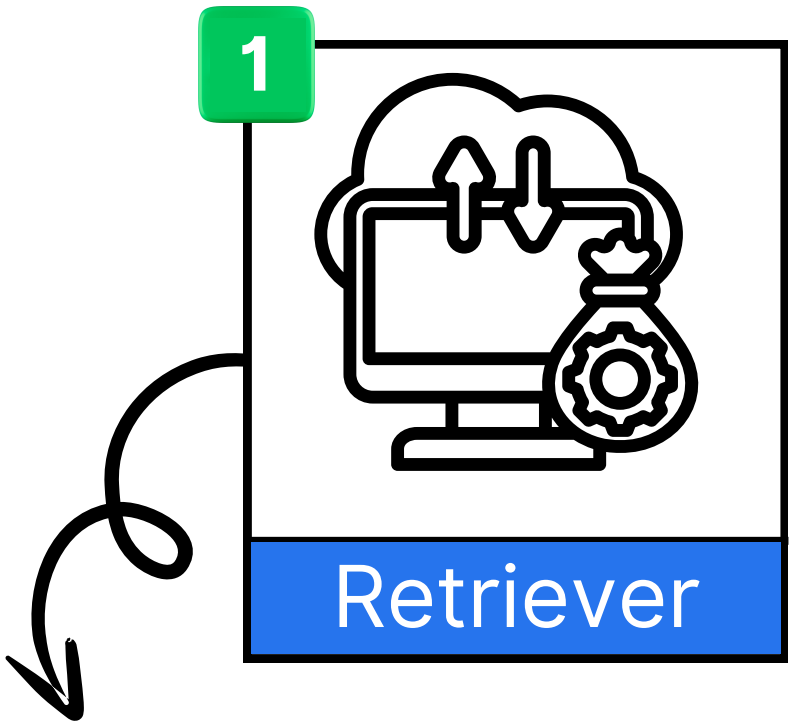

The retriever searches for and finds the best information from databases or websites, like a search engine that quickly pulls up relevant details.

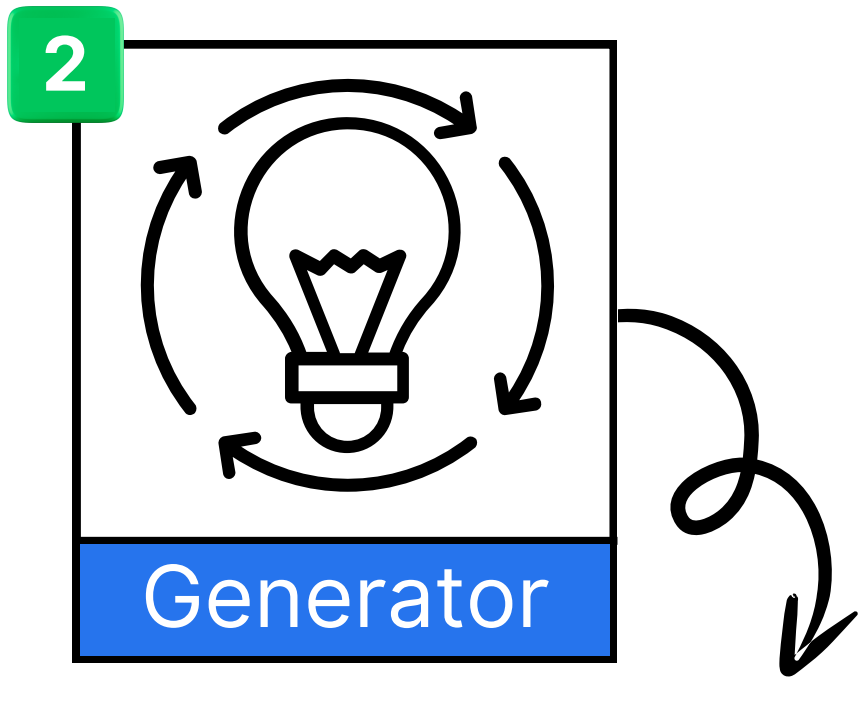

The generator takes that information and combines it with what it already knows to create a clear, human-like response.

Working of a RAG

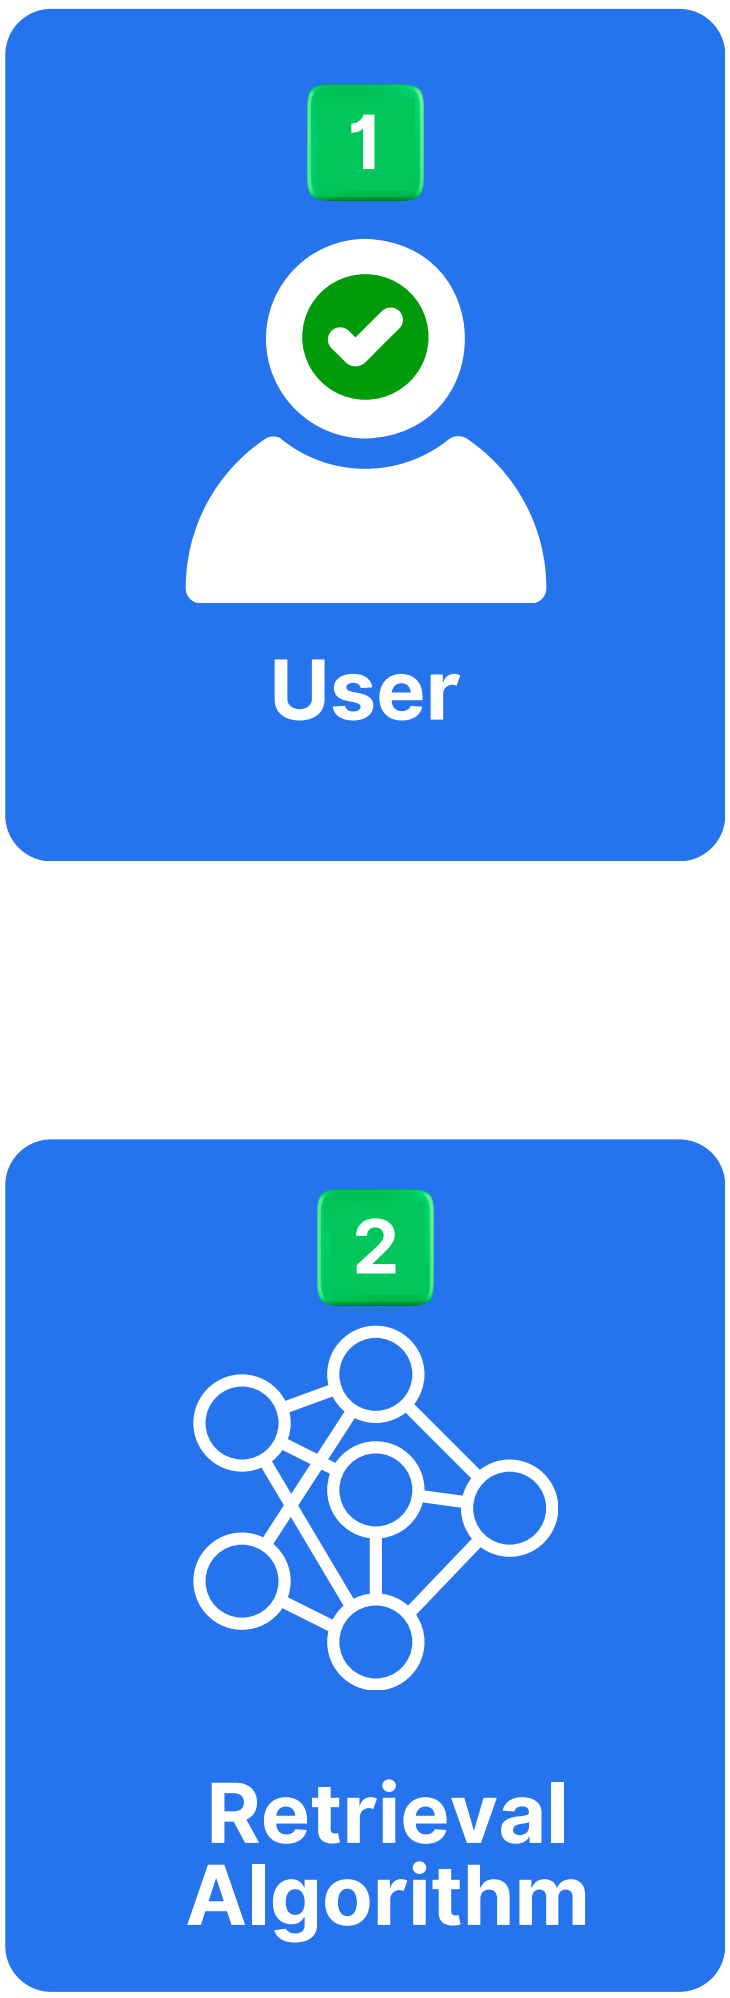

In the first step, the user ask the query to the LLM.

The query is then sent to the **Retrieval algorithm**, that is responsible to fetch the relevant documents from the knowledge base.

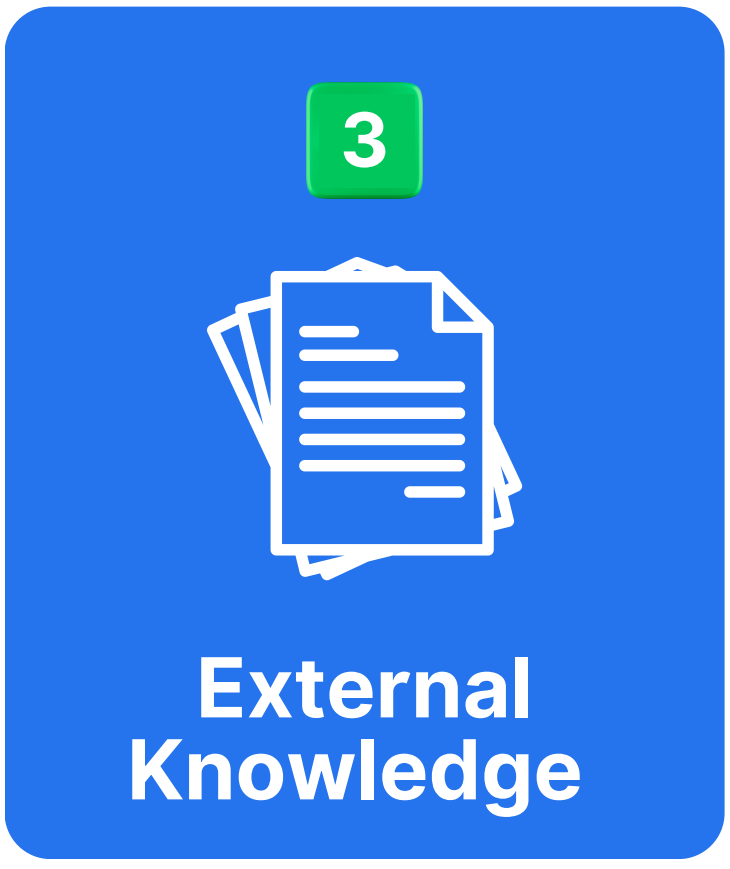

This **External knowledge base** is the source from where the Retrieval algorithm is fetching the relevant documents.

The retrieved documents, along with the original query, are sent to the language model (LLM).

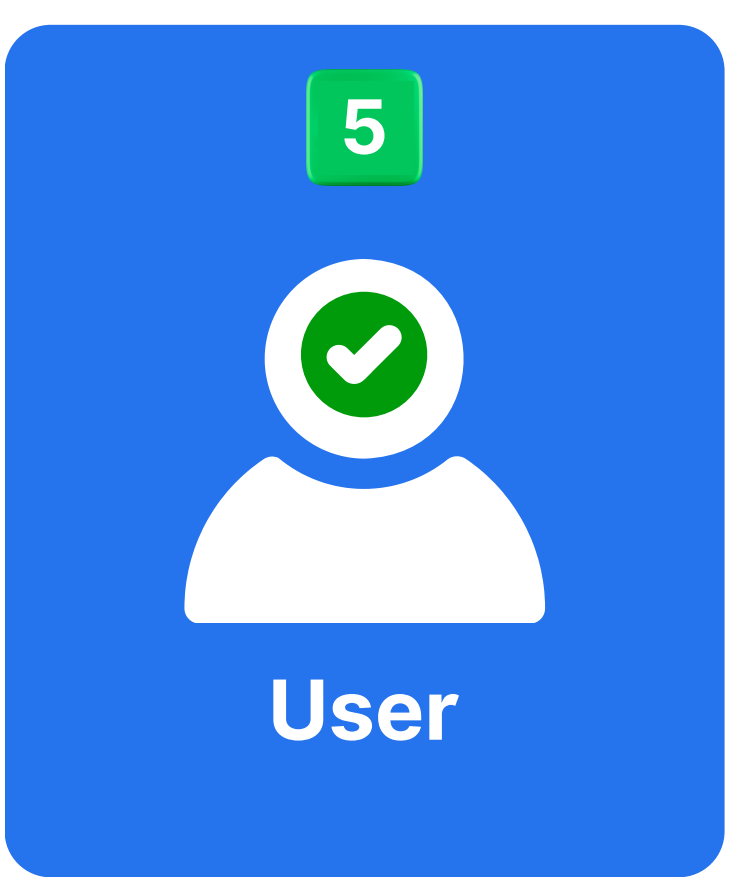

The generator processes both the query and the relevant documents to generate a response, which is then sent back to the user.

Analytics Vidhya

FREE COURSE

RAG

We are offering a FREE Course on “Building Your first RAG System using LlamaIndex”