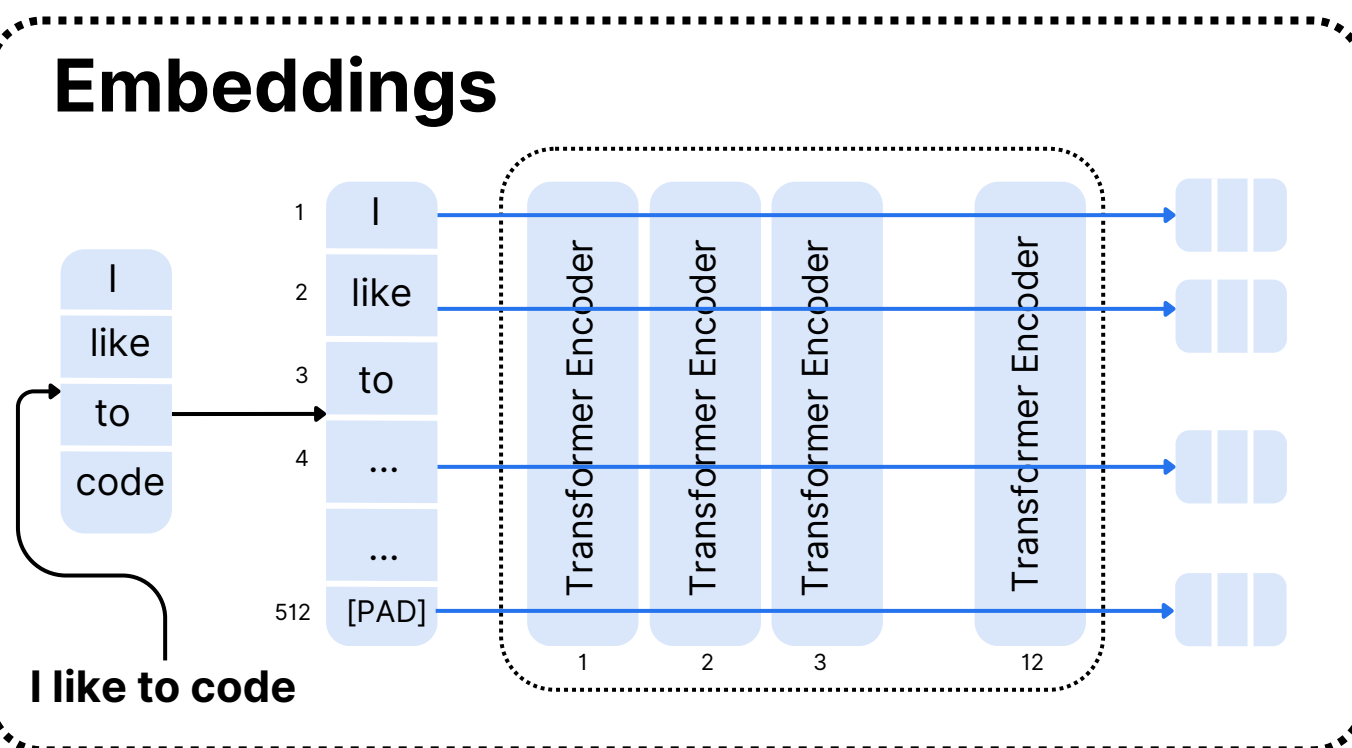
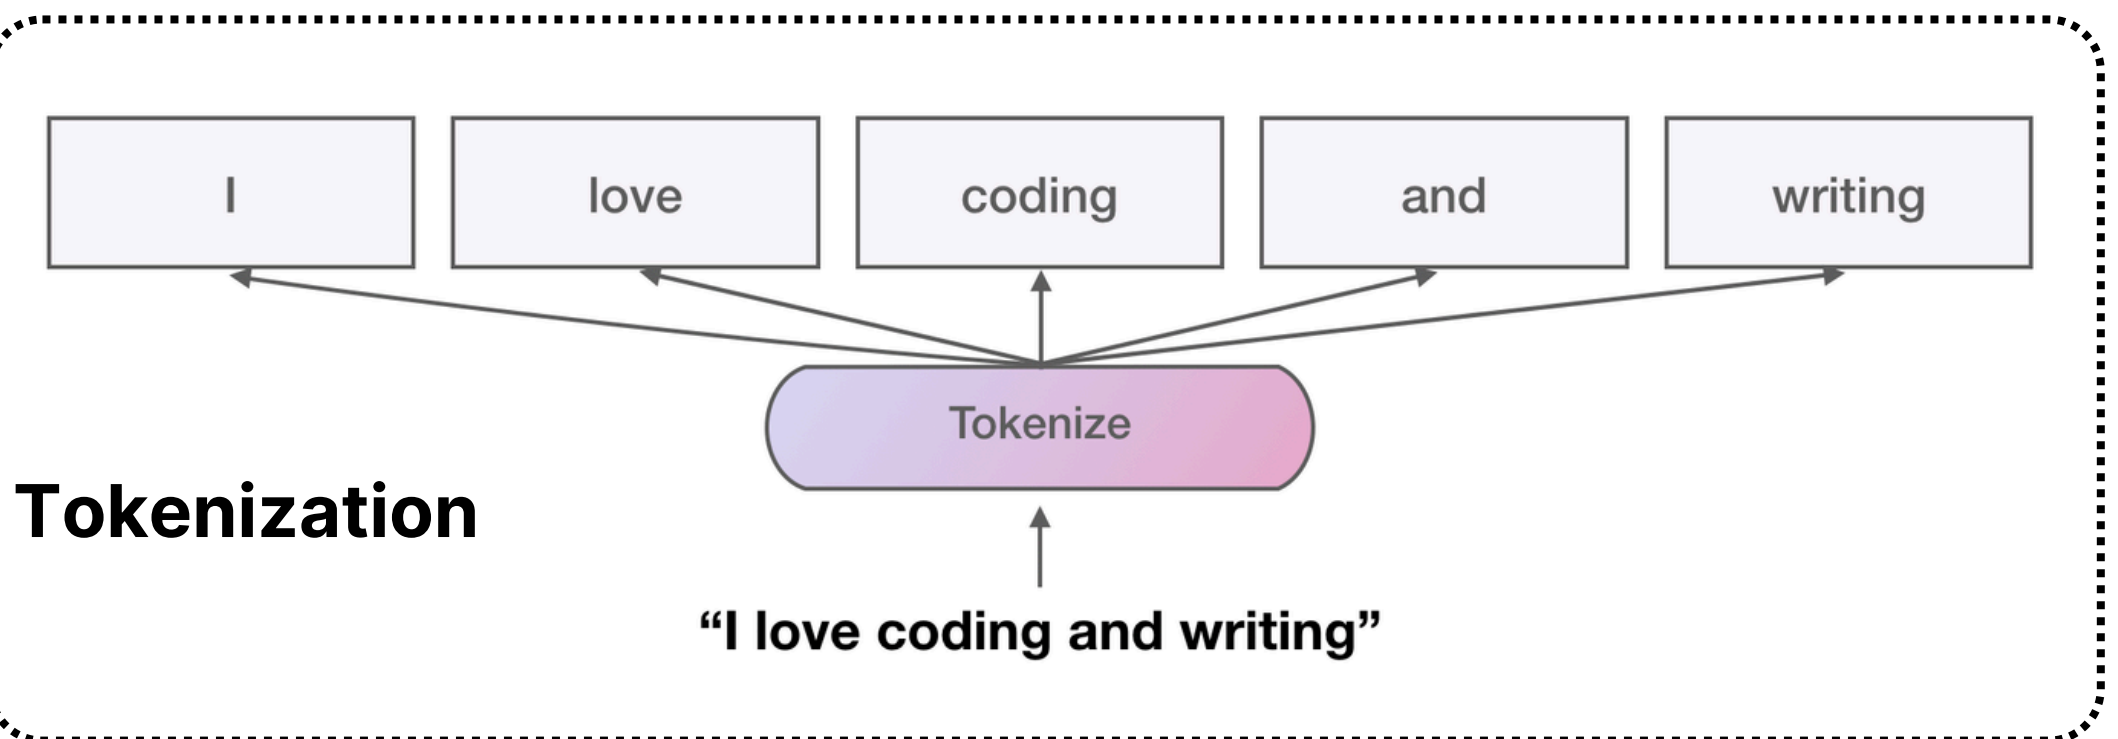


Mastering LLMs

Day 5: Tokenization and Embeddings



```

# Importing the necessary library
from transformers import AutoTokenizer

# Load a pre-trained tokenizer (e.g., BERT tokenizer)
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")

# Input text to tokenize
text = "Transformers are revolutionizing natural language processing."

# Tokenize the text
tokenized_output = tokenizer(text)

# Print the raw tokenized output
print("Raw Tokenized Output:")
print(tokenized_output)

# Print each component of the tokenized output
print("\nTokens (Text):",
      tokenizer.convert_ids_to_tokens(tokenized_output['input_ids']))
print("Input IDs:", tokenized_output['input_ids'])

print("Attention Mask:", tokenized_output['attention_mask'])

# Decode the token IDs back to text
decoded_text = tokenizer.decode(tokenized_output['input_ids'])
print("\nDecoded Text:", decoded_text)

# Additional exploration: Tokenization with padding and truncation
tokenized_with_padding = tokenizer(
    text,
    padding="max_length", # Adds padding
    max_length=12, # Ensures the sequence length is at most 12
    truncation=True # Truncates longer sequences
)

print("\nTokenized with Padding and Truncation:")
print("Input IDs:", tokenized_with_padding['input_ids'])
print("Attention Mask:", tokenized_with_padding['attention_mask'])

# Batch tokenization
batch_texts = ["Transformers are amazing.", "AI is the future."]
batch_tokenized = tokenizer(batch_texts, padding=True, truncation=True)

print("\nBatch Tokenized Output:")
for i, text in enumerate(batch_texts):
    print(f"Text {i + 1}:")
    print("Input IDs:", batch_tokenized["input_ids"][i])
    print("Attention Mask:", batch_tokenized["attention_mask"][i])
  
```

```

from transformers import AutoTokenizer, AutoModel
import torch

# Load pre-trained tokenizer and model
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")
model = AutoModel.from_pretrained("bert-base-uncased")

# Input text
text = "Transformers are revolutionizing AI and NLP."

# Tokenize the input text
tokenized_output = tokenizer(text, return_tensors="pt", padding=True, truncation=True)

# Extract input IDs and attention mask
input_ids = tokenized_output['input_ids'] # Shape: [batch_size, sequence_length]
attention_mask = tokenized_output['attention_mask'] # Shape: [batch_size, sequence_length]

# Pass the tokenized input through the model to get embeddings
with torch.no_grad():
    model_output = model(input_ids=input_ids, attention_mask=attention_mask)


# Extract the embeddings (last hidden states)
embeddings = model_output.last_hidden_state # Shape: [batch_size, sequence_length, embedding_dim]
embedding_dim

# Display embedding details
print("Embeddings Shape:", embeddings.shape)
print("First Token Embedding:", embeddings[0][0]) # Embedding for the [CLS] token

# Convert token IDs back to tokens for clarity
tokens = tokenizer.convert_ids_to_tokens(input_ids[0])
print("\nTokens and their embeddings:")
for token, embedding in zip(tokens, embeddings[0]):
    print(f"Token: {token}, Embedding (first 5 dimensions): {embedding[:5]}")
  
```

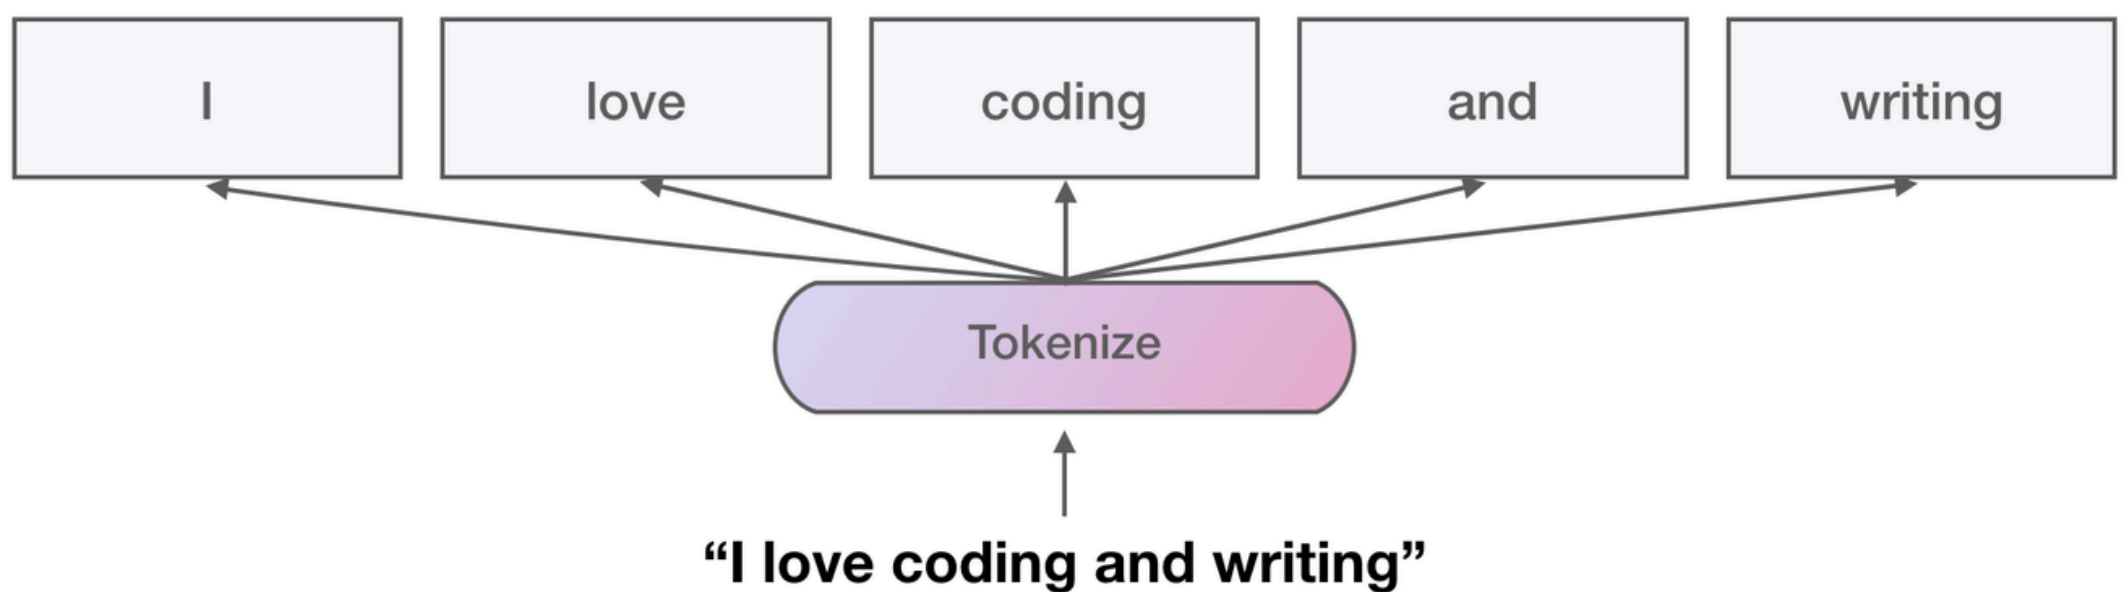
Embeddings

Tokenization



Tokenization

- Tokenization is the process of breaking down text into smaller units, called tokens, that a deep learning model can process.




Types of Tokenization

- Word-based Tokenization
- Subword-based Tokenization
- Character-based Tokenization

Word-based Tokenization

- Each word is treated as a token.



- **Limitation:** Cannot handle unseen words or word variants effectively (e.g., "transformer" vs. "transformers").

Subword-based Tokenization

- Breaks words into smaller **meaningful units**, like **prefixes**, **suffixes**, or **roots**.
- Common algorithms:
 - Byte Pair Encoding (BPE)
 - WordPiece
 - SentencePiece
- Example: "transformers" → ["transform", "##ers"] (WordPiece) or ["trans", "former", "s"] (BPE).

Character-based Tokenization

- Each **character** is treated as a token.
- Example: "Hi" → ["H", "i"]
- Advantage: Handles any input but can result in long sequences.

Here is a Python code to demonstrate the working of **tokenization** in transformers using the 🤗 **Hugging Face**

```
# Importing the necessary library
from transformers import AutoTokenizer

# Load a pre-trained tokenizer (e.g., BERT tokenizer)
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")

# Input text to tokenize
text = "Transformers are revolutionizing natural language processing."

# Tokenize the text
tokenized_output = tokenizer(text)

# Print the raw tokenized output
print("Raw Tokenized Output:")
print(tokenized_output)

# Print each component of the tokenized output
print("\nTokens (Text):",
      tokenizer.convert_ids_to_tokens(tokenized_output['input_ids']))
print("Input IDs:", tokenized_output['input_ids'])

print("Attention Mask:", tokenized_output['attention_mask'])

# Decode the token IDs back to text
decoded_text = tokenizer.decode(tokenized_output["input_ids"])
print("\nDecoded Text:", decoded_text)

# Additional exploration: Tokenization with padding and truncation
tokenized_with_padding = tokenizer(
    text,
    padding="max_length", # Adds padding
    max_length=12,        # Ensures the sequence length is at most 12
    truncation=True        # Truncates longer sequences
)

print("\nTokenized with Padding and Truncation:")
print("Input IDs:", tokenized_with_padding["input_ids"])
print("Attention Mask:", tokenized_with_padding["attention_mask"])

# Batch tokenization
batch_texts = ["Transformers are amazing.", "AI is the future."]
batch_tokenized = tokenizer(batch_texts, padding=True, truncation=True)

print("\nBatch Tokenized Output:")
for i, text in enumerate(batch_texts):
    print(f"Text {i + 1}:")
    print("  Input IDs:", batch_tokenized["input_ids"][i])
    print("  Attention Mask:", batch_tokenized["attention_mask"][i])
```

Download
the **PDF**
for a
closer
look at
the code

The explanation of
this code is
provided on the
next slide →

How This Code Works

Tokenizer Initialization:

- A pre-trained tokenizer (bert-base-uncased) is loaded, which matches the vocabulary of the BERT model.

Tokenization:

- The input text is tokenized into IDs, and an attention mask is generated.
- Special tokens like [CLS] (start of sequence) and [SEP] (end of sequence) are automatically added.

Decoding:

- The tokenized IDs are converted back into readable text to see how the tokenizer represents it.

Padding and Truncation:

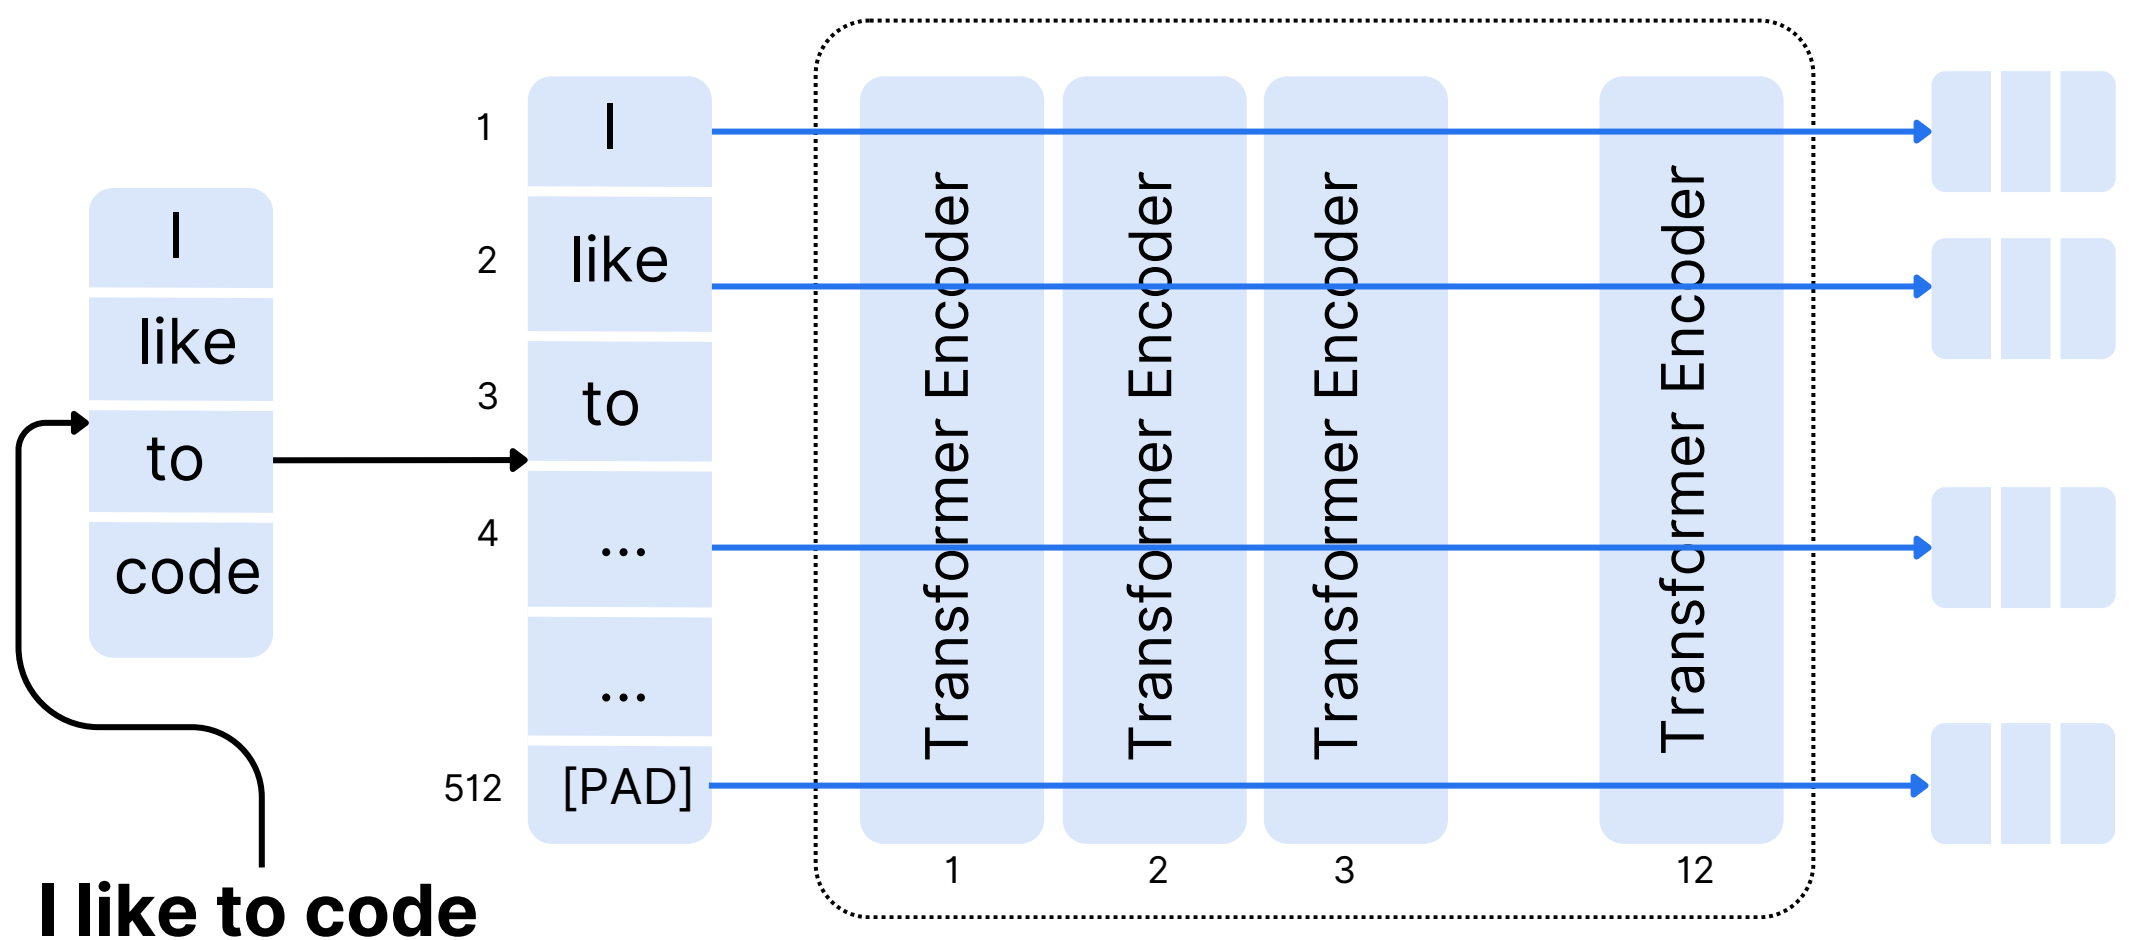
- Ensures that tokenized sequences conform to a fixed length by adding padding tokens or truncating longer texts.

Batch Tokenization:

- Tokenizes multiple texts at once, ensuring consistent sequence lengths via padding.

Embeddings

- Embeddings are dense numerical vector representations of tokens, capturing their meanings, relationships, and context in a lower-dimensional space.



How It Works:

- Each token is mapped to a dense vector using an embedding layer (usually implemented as a lookup table).
- The vectors are learned during training to capture semantic relationships.

Features of Embeddings

Context-Free Embeddings

- Each word has a fixed representation, regardless of context.
- Example: Word2Vec, GloVe.

Contextualized Embeddings

- Representations change based on context, thanks to transformer architecture.
- Example: "bank" in "river bank" vs. "financial bank" will have different embeddings.
- Achieved using mechanisms like self-attention.

Mathematical Representation

- Suppose a vocabulary has V tokens and embedding size is d . The embedding layer is a matrix E of size $V \times d$ times $V \times d$.
- For a token with index i , its embedding vector is $E[i]$.

Why Important?

Embeddings enable transformers to:

- Capture the semantic similarity between words (e.g., "king" and "queen").
- Represent syntactic and grammatical information.
- Work effectively with numeric data instead of raw text.

Interaction in Transformers

- **Tokenization** prepares input text, converting it to token IDs.
- **Embedding Layer maps** these IDs to vectors.
- These vectors are passed through the transformer's layers for processing.

Example Pipeline

- Input: "I love AI"
- **Tokenization:** [101, 1045, 2293, 2143, 102] (using WordPiece with special tokens)
- **Embedding:** Vectors corresponding to these IDs, ready for transformer layers.

Here is a Python code to demonstrate the working of embedding in transformers:

```
from transformers import AutoTokenizer, AutoModel
import torch

# Load pre-trained tokenizer and model
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")
model = AutoModel.from_pretrained("bert-base-uncased")

# Input text
text = "Transformers are revolutionizing AI and NLP."

# Tokenize the input text
tokenized_output = tokenizer(text, return_tensors="pt", padding=True, truncation=True)

# Extract input IDs and attention mask
input_ids = tokenized_output['input_ids'] # Shape: [batch_size, sequence_length]
attention_mask = tokenized_output['attention_mask'] # Shape: [batch_size, sequence_length]

# Pass the tokenized input through the model to get embeddings
with torch.no_grad():
    model_output = model(input_ids=input_ids, attention_mask=attention_mask)

# Extract the embeddings (last hidden states)
embeddings = model_output.last_hidden_state # Shape: [batch_size, sequence_length, embedding_dim]
# Display embedding details
print("Embeddings Shape:", embeddings.shape)
print("First Token Embedding:", embeddings[0][0]) # Embedding for the [CLS] token

# Convert token IDs back to tokens for clarity
tokens = tokenizer.convert_ids_to_tokens(input_ids[0])
print("\nTokens and their embeddings:")
for token, embedding in zip(tokens, embeddings[0]):
    print(f"Token: {token}, Embedding (first 5 dimensions): {embedding[:5]}")
```

Download
the **PDF**
for a
closer
look at
the code

The
explanation of
this code is
provided on
the next slide



How This Code Works

Tokenizer and Model:

- **AutoTokenizer**: Loads a tokenizer compatible with bert-base-uncased.
- **AutoModel**: Loads the BERT model to compute embeddings.

Tokenization:

- **tokenizer**: Converts input text into input_ids and attention_mask required for the model.
- **padding=True**: Ensures all sequences are padded to the same length.
- **truncation=True**: Truncates sequences that exceed the model's maximum length.

Model Forward Pass:

- The tokenized input is passed through the model using model(input_ids, attention_mask).
- The output includes:
 - **last_hidden_state**: Token embeddings for each position in the input sequence.
 - Other optional outputs (not used here).

Embedding Extraction:

- last_hidden_state contains embeddings for all tokens in the input sequence, including special tokens like [CLS] and [SEP].

Token Inspection:

- Tokens are paired with their corresponding embeddings for inspection.
- Only the first 5 dimensions of each embedding vector are displayed for readability.

January

2025

Sunday	Monday	Tuesday	Wednesday	Thursday	Friday	Saturday
			1	2	3	4
5	6 ✓	7 ✓	8 ✓	9 ✓	10 ✓	11
12	13	14	15	16	17	18
19	20	21	22	23	24	25
26	27	28	29	30	31	

Stay Tuned for **Day 6** of

Mastering LLMs